SYLVR
RISKS to RESOLVE to REWARDS

SYLVR | Assignment | NLP & Agentic AI Intern

Challenge: Build a **Conversational Database Agent** that can have conversation with user about data stored in a database by

1. Mapping natural language text into various query types (definitions, filters, aggregation, trends, comparisons, etc.)
2. Executing appropriate database queries on relevant collections to obtain relevant information, or else, flag gaps in data
3. Providing quality responses and handling follow up text to maintain conversation context across multiple exchanges
4. Extracting objective, actionable insights on database gaps and user's emotional tone, query intent for World Model

Key Requirements

1. **Database integration** – Use a real MongoDB database; no PDF/document querying

1.1 sample_analytics database (finance domain) from MongoDB's free Atlas cluster is sufficient for base requirements

1.2 For extra credit go beyond default database by using a more complex financial database meeting these requirements:

1.2.1 More collections, Financial numeric data fields, Nested fields, One or more high volume collections (>500 documents)

2. **Query execution** – Handle various question types and schema, metadata, vector embeddings/ similarity for relevance
3. **Conversation context** – Handle follow-up text referencing prior conversation parts and 'memorize' multiple exchanges
4. **Deliverables** – GitHub repo with working code, README.md with setup and brief architecture explanation, dataset used (upload and give link in README file) and short demo video demonstrating various user questions

4.1 **NOTE: If you can't complete Secondary and Bonus features fully in given time then** include a brief writeup in your README explaining how you would implement the remaining, and what further you would implement given more time. A thoughtful writeup about features you didn't have time to build will be duly considered.

What we are evaluating

PRIMARY (Must Have)	**SECONDARY (Nice to Have)**	**BONUS (Optional Extra Credit)**
Database connectivity and query execution on atleast 3 collections to simulate real-world query	Agent's Context use for conversation management with simple in-memory data structures or LangChain's memory modules	World Model insights dashboard
Natural language understanding and query translation	Error handling for 'ambiguous/ impossible' queries	Advanced agentic reasoning
Basic conversation memory and context handling	Multiple database operations in one conversation	Query optimization
Working solution in Colab/ Jupyter with end-to-end demo	UI using any simple framework like Gradio, Streamlit, React, etc.	Creative solutions to complex problems
Database - MongoDB	Clean, organized code architecture	Voice interface (STT/ TTS)
Programming language – Python (backend, logic & DB interactions)		Docker file to containerize your application for easy deployment

NOTE: We are primarily evaluating AI/backend functionality of Query, Response, Context quality with clean architecture & code, a working implementation and problem-solving approach with technical planning. Focus on making the Agent smart.